Strategic Optimization of Continuous Sign Language Recognition Frameworks for the PHOENIX-2014-T Benchmark

The field of Continuous Sign Language Recognition (CSLR) has undergone a significant paradigm shift between 2023 and 2025, moving away from simple sequence-to-sequence mapping toward nuanced spatio-temporal modeling and complex alignment supervision.[1] The RWTH-PHOENIX-Weather 2014-T dataset remains the gold standard for evaluating German Sign Language (Deutsche Gebärdensprache, DGS) systems, representing a uniquely challenging environment due to its limited sample size of approximately 7,000 videos and the intricate nature of weather-specific vocabulary.[3] Achieving a competitive Word Error Rate (WER) on this benchmark requires more than a deeper convolutional backbone; it necessitates a sophisticated integration of visual alignment, temporal resolution management, and domain-specific pre-training.[5]

Theoretical Foundations of Visual Alignment and CTC Stability

The primary challenge in CSLR is the weak supervision inherent in sentence-level annotations, where the precise temporal boundaries of individual glosses are unknown.[1] Connectionist Temporal Classification (CTC) has traditionally addressed this by allowing for alignment-free training, but its reliance on the "blank" label often leads to a degenerate solution known as CTC collapse.[8] When an auxiliary Visual Alignment Constraint (VAC) is introduced to regularize the feature extractor, the model frequently over-optimizes for the blank prediction, resulting in a WER that approaches 100%.[8]

Mechanics of the Visual Alignment Constraint and Mitigation of Collapse

The failure of VAC in standard hybrid architectures often stems from a gradient dominance problem where the auxiliary CTC head provides a simpler path to minimize loss than the primary sequence decoder.[10] In the context of a ResNet-18 or ResNet-50 backbone, applying an auxiliary CTC loss directly to the spatial features before the temporal modeling stage forces the CNN to become highly discriminative before the temporal sequence has been properly structured.[8] This conflict is exacerbated when the auxiliary head is not sufficiently decoupled from the main feature extractor.[8]

Recent research suggests that a projection layer, typically a linear bottleneck followed by Batch Normalization and ReLU, is essential between the convolutional backbone and the auxiliary

CTC head.[8] This projection layer serves as a buffer, allowing the auxiliary head to learn its own internal representation of the gloss sequence without directly backpropagating catastrophic interference into the spatial features.[8] Furthermore, the weighting of this auxiliary loss, often denoted as λ, must be carefully managed. A static weight of 1.0 is frequently too aggressive for the limited data in PHOENIX-2014-T. Implementing a dynamic weight schedule—starting at $\lambda = 0.0$ and gradually warming up to $\lambda = 0.1$ over the first 10,000 iterations—allows the spatial backbone to stabilize its ImageNet-derived weights before adapting to the monotonic constraints of DGS.[10]

Loss Component	Optimization Target	Role in Gradient Flow	Typical Weight Range
Primary CTC	Sequence Alignment	Global temporal structure	1.0 (baseline) [11]
Cross-Entropy	Linguistic Context	Sentence-level transition modeling	0.5 - 2.0 [5]
VAC (Auxiliary)	Feature Discriminability	Backbone regularization	0.01 - 0.2 [10]
Distillation Loss	Alignment Consistency	Frame-wise consistency	0.1 - 0.5 [12]

The "spiky problem" inherent in CTC outputs, where predictions form sharp spikes separated by long durations of high-confidence blanks, is another driver of instability.[9] By reinterpreting CTC training as an iterative fitting process, researchers have identified that the dominance of blank labels is a byproduct of the model attempting to maximize the probability of all possible paths.[9] To counteract this, recent frameworks incorporate a frame-wise regularization term, which forces the model to distribute its probability mass more evenly across the duration of a sign, thereby providing the Attention decoder with a richer and more continuous feature stream.[9]

Memory-Efficient Spatio-Temporal Modeling

The failure of 3D CNN backbones like R3D-18 in many CSLR pipelines is rarely due to the 3D kernels themselves but rather the memory-driven compromises they necessitate.[2] On a 16GB VRAM GPU, the use of R3D-18 often requires limiting input sequences to fewer than 64 frames or utilizing aggressive temporal downsampling with a stride of 4 or 8.[2] Because DGS signs are often articulated at high speeds—sometimes spanning as few as 10 frames—this downsampling effectively "erases" the gesture from the input sequence before it reaches the classifier.[15]

Alternatives to 3D Convolutions: TSM and TLP

To capture motion without the parameter explosion of 3D kernels, the Temporal Shift Module (TSM) and Temporal Lift Pooling (TLP) have emerged as superior alternatives for memory-constrained environments.[6] TSM facilitates inter-frame communication by shifting a portion of the feature channels along the temporal dimension, effectively enabling a 2D CNN to perceive motion at zero additional computational cost.[6]

Temporal Lift Pooling (TLP) represents a more sophisticated approach derived from the Lifting Scheme in signal processing.[17] Unlike traditional max or average pooling, TLP decomposes the temporal signal into different sub-bands, preserving high-frequency motion patterns that are typically lost during downsampling.[17]

Technique	Mechanism	Computational Overhead	Key Benefit for CSLR
3D CNN (R3D-18)	Spatio-temporal kernels	High VRAM / High Ops	Holistic motion capture [2]
TSM	Channel shifting	Zero	Efficient short-term dynamics [6]
TLP	Sub-band decomposition	Low	Preserves high-frequency motion [17]
1D CNN Stack	Temporal convolution	Low	Large temporal receptive field [18]

For a single-GPU setup with 16GB VRAM, the most effective architecture for PHOENIX-2014-T involves a ResNet-18 backbone integrated with TSM in every other residual block.[6] This maintains the high temporal resolution of the input (25 FPS) while allowing the network to learn the relative motion of the hands and face without the aggressive downsampling required by R3D architectures.[6]

Adaptive Temporal Resolution Management

A recurring issue in the baseline architecture is the use of fixed-stride downsampling, which fails to account for the varying speeds of different signers.[2] The Dual-stage Temporal Perception Network (DTPNet) and the CNN Transformer with Adaptive Temporal Hierarchical Attention

(CT-ATHA) address this through learnable temporal mechanisms.[2]

Adaptive Temporal Pooling (ATP) uses the temporal entropy of each frame to determine its importance.[2] High-entropy frames, which represent rapid handshape transitions or significant movements, are preserved, while low-entropy "static" frames are compressed.[2] This mechanism acts as a dynamic stride, allowing the model to focus its 16GB VRAM budget on the most informative segments of the video.[2] Complementing this, Learnable Temporal Gates (LTG) act as adaptive filters that prune transitional and rest periods, represented mathematically as $F_{gated} = G(F_{in}) \odot F_{in}$, where G is a sigmoid-activated learned function.[2] This pruning ensures that the expensive Transformer layers are only processing meaningful gestural content, significantly reducing the likelihood of overfitting to background noise.[2]

Architectural Refinements for Local and Global Context

The 6-layer Transformer Encoder used in the baseline model possesses a global receptive field but lacks the inductive bias for the local temporal correlations essential for sign recognition.[5] In sign language, the transition between handshapes (local dynamics) is as important as the sentence-level structure (global dynamics).[2]

The Conformer: Integrating Local Connectivity

The Conformer architecture has recently been adapted for vision-based tasks like CSLR under the name ConSignformer.[5] By integrating a convolution module into the standard Transformer block, the Conformer captures local temporal features (similar to an n-gram in speech) while the self-attention mechanism handles long-range dependencies.[5] For PHOENIX-2014-T, where co-articulation can blur the boundaries between signs, the Conformer's ability to maintain high-resolution local features is vital for reducing the confusion between visually similar glosses.[5]

Furthermore, the introduction of Cross-Modal Relative Attention (CMRA) allows the model to better utilize the relationships between different spatial regions, such as the relative positioning of the hand and the face.[5] This is particularly relevant for DGS, where the location of a sign on the body can fundamentally alter its meaning.[5]

Light-weight Hand and Face Attention

Capturing hand and facial nuances without the memory cost of a multi-stream pose-estimation network is a key requirement for consumer-grade GPUs.[20] The Motor Attention Module (MAM) offers a lightweight solution by utilizing frame-differencing as a motion prior.[2] MAM weights the spatial features generated by the CNN backbone, forcing the model to attend to regions of high motion (typically the hands) and suppress the static background.[2]

In a similar vein, the Correlation Network (CorrNet+) uses "Correlation Modules" to compute the cross-spacetime correlations in local neighborhoods.[6] Rather than processing frames as independent static images, CorrNet+ constructs human body trajectories by identifying how spatial patches move across consecutive frames.[6] This trajectory-centric approach is highly effective for PHOENIX-2014-T because it mimics the way human interpreters perceive signs—not as individual frames, but as continuous paths of motion.[6]

Attention Mechanism	Input Requirement	Memory Impact	Domain Focus
Standard Attention	Global features	High ($O(T^2)$)	Global context [5]
MAM (Motor)	Frame difference	Low	Active motion regions [2]
Correlation Module	Adjacent patches	Medium	Body trajectories [6]
Pose Heatmaps	Off-the-shelf detector	Very High	Skeleton-centric [20]

Self-Supervised Sign Representation Learning

The small size of PHOENIX-2014-T often leads to severe overfitting, as evidenced by the user's CorrNet+ results (32% train WER vs. 79% val WER).[7] Self-supervised pre-training on unlabeled sign data or through masked modeling is the most effective way to improve generalization.[23]

Masked Autoencoding with Sign Priors (SignRep)

The SignRep framework (2025) introduces a scalable method for learning sign representations without the need for sentence-aligned labels.[24] It uses a Masked Autoencoder (MAE) approach, but rather than reconstructing raw pixels, it reconstructs "sign priors"—geometric cues that capture the essence of signing.[24]

1. **Keypoint Priors**: Reconstructing the 3D coordinates of hand and body joints forces the model to learn the spatial kinematics of DGS.[24]
2. **Joint Angle Priors**: Reconstructing the flexion and extension of fingers helps the model distinguish between subtle handshapes like the "O" and "C" shapes common in weather terminology.[24]

3. **Distance Priors**: By predicting the distance between fingertips and the face, the model learns the "signing space," which is critical for understanding signs that involve contact with the body.[24]
4. **Signer Activity**: A binary prior indicating whether a hand is in an active or resting state helps the model ignore non-lexical transitions.[24]

Crucially, these priors are only used during the pre-training phase. At inference, only the RGB backbone is required, making the model highly efficient for 16GB VRAM deployment while benefiting from the robust features learned during the multi-modal pre-training stage.[24]

SignBERT+ and Masked Clip Modeling

SignBERT+ extends this concept by treating hand poses as "visual tokens" in a frame-by-frame manner.[23] It employs multi-level masking strategies—joint, frame, and clip levels—to mimic common failures in real-world signing, such as motion blur and occlusion.[23]

- **Clip-Level Masking**: Masking consecutive frames (typically 8 frames) forces the model to capture the temporal dynamics and flow of a sign sequence.[23]
- **Model-Aware Hand Prior**: By incorporating the MANO statistical hand model into the decoder, SignBERT+ ensures that its reconstructed poses are kinematically plausible, preventing the model from learning "impossible" hand configurations.[23]

For an RTX 5070 Ti user, the best path forward is to utilize a backbone pre-trained with SignBERT+ or SignRep. This "sign-aware" initialization is far more effective than ImageNet pre-training, as it bridges the domain gap between natural images and the specialized gestures of DGS.[24]

Advanced Loss Functions and Knowledge Distillation

The standard CTC loss treats all paths equally, which often ignores the fact that some temporal alignments are more linguistically plausible than others.[12] To resolve this, the Distillation Connectionist Temporal Classification (DCTC) loss has been introduced.[12]

DCTC and Latent Alignment Estimation

DCTC addresses the "alignment inconsistency" problem by using a Maximum-A-Posteriori (MAP) estimate of the latent alignment to regularize the frame-wise predictions.[12] This acts as a form of self-distillation, where the model's most confident alignment path is used to supervise the less certain frames.[13] This technique is particularly effective for PHOENIX-2014-T, as it provides a more stable gradient than standard CTC when the training data is noisy or limited.[12]

Another variant, Focal CTC, modulates the importance of different glosses during training.[13] In weather forecasts, certain signs (e.g., "RAIN," "SUN") are extremely common, while specific locations or rare weather patterns are infrequent. Focal CTC prevents the model from ignoring these rare signs by increasing their relative loss weight, helping to lower the overall WER by improving performance on the long tail of the vocabulary.[13]

Cross-Resolution and Multi-Modal Knowledge Distillation

When hardware constraints prevent the use of high-resolution images or deep 3D networks, Knowledge Distillation (KD) provides a way to transfer performance from a "Teacher" model to a "Student" model.[21]

Cross-Resolution Knowledge Distillation (CRKD) is particularly relevant for the RTX 5070 Ti.[27] A teacher model is trained on high-resolution (224x224) frames to capture fine-grained hand details.[28] Once trained, its weights are locked, and it supervises a student model (LRINet) receiving low-resolution (112x112) inputs.[27] CRKD ensures that the low-resolution features generated by the student retain the semantic richness of the high-resolution teacher, allowing for a 1.74x increase in inference speed and a significant reduction in VRAM usage without a corresponding drop in accuracy.[27]

Hierarchical Knowledge Distillation (SignKD) transfers knowledge from multi-modal sources (RGB + Keypoints + Optical Flow) to a single-modal RGB model.[21] This allows the final CSLR model to benefit from motion and skeletal information during training while only requiring RGB video at test time.[21]

Distillation Type	Teacher Input	Student Input	Key Objective
CRKD	High-Res (224x224)	Low-Res (112x112)	Memory efficiency [28]
SignKD	Multi-modal (Flow/Pose)	RGB-only	Motion awareness [21]
Self-Distillation (DCTC)	Best MAP path	All paths	Alignment stability [12]
Cross-Lingual	Multilingual Dictionary	Target DGS Data	Data scarcity mitigation [7]

Data Scarcity and Cross-Lingual Augmentation

The PHOENIX-2014-T corpus is limited, but sign languages often share universal visual signals.[7] Cross-lingual signs—gestures that appear in different sign languages but convey similar meanings—can be used as auxiliary training data.[7]

Sign Dictionary Construction and Mapping

A 2023 ICCV paper demonstrates that identifying CSL-to-DGS (Chinese to German) sign mappings can significantly improve CSLR performance.[7] By building isolated sign dictionaries and identifying visually similar signs, researchers can leverage the much larger CSL-Daily dataset to augment the PHOENIX-2014-T training source.[3] This approach treats CSLR as a transfer learning problem, using the multilingual corpus to build a more robust visual representation that is less prone to the "forgetting" seen in small datasets.[3]

Furthermore, "Sign Augmentation" can be achieved by using a SOTA model like TwoStream-SLR as a segmentor.[16] The model identifies the approximate boundaries of signs in a continuous video, allowing for the extraction of "pseudo-isolated" signs.[16] These segments are then cropped and added back to the training set with their corresponding gloss labels, effectively tripling the number of training samples and providing the model with more examples of co-articulation at sign boundaries.[16]

Synthesis of Findings and Actionable Optimization Roadmap

For a CSLR practitioner utilizing a 16GB RTX 5070 Ti on PHOENIX-2014-T, the path from a 50.91% WER to the SOTA 18-20% range requires a phased approach that addresses the failures of VAC, memory-intensive backbones, and limited training data.

Phase 1: Backbone and Visual Alignment Stabilization

The most immediate cause of failure in current experiments is the CTC collapse induced by the VAC loss. To resolve this:

- **Introduce an Auxiliary Head**: Add a linear bottleneck layer (projection head) between the CNN and the auxiliary CTC loss.[8]
- **Implement Loss Weight Warm-up**: Start the VAC weight at $\lambda = 0.0$ and linearly increase it to $\lambda = 0.1$ over the first 5 epochs.[10]
- **Replace ResNet-50 with ResNet-18 + TSM**: ResNet-50 offers more depth, but the memory savings from ResNet-18 allow for a higher frame rate and the inclusion of the Temporal Shift Module, which is more critical for CSLR than static depth.[6]

Phase 2: Temporal Modeling and Adaptive Resolution

Once the backbone is stable, the 6-layer Transformer Encoder should be refined:

- **Adopt the Conformer Structure**: Replace the Transformer layers with Conformer blocks to better capture local handshape movements.[5]

- **Integrate Adaptive Temporal Pooling**: Use a temporal entropy-based pooling layer to reduce the 120-frame sequences to 30 frames before the Conformer layers, enabling larger batch sizes.[2]
- **Apply Motor Attention**: Add the lightweight MAM to the ResNet-18 stages to focus the model on the signer's hands without the need for skeletal detection.[2]

Phase 3: Domain-Specific Pre-training and Advanced Alignment

To overcome the data scarcity of PHOENIX:

- **Sign-Specific SSP**: Pre-train the ResNet-18 + TSM backbone on the PHOENIX-2014-T videos using masked frame prediction.[23] If possible, use the SignRep approach of predicting fingertip distances and joint angles from a 3D pose-estimation teacher.[24]
- **Switch to DCTC Loss**: Replace the standard CTC loss with DCTC to enforce frame-wise alignment consistency and prevent the "spiky problem".[12]
- **Augment via Segmentation**: Use a pre-trained TwoStream-SLR model to segment the PHOENIX videos into isolated signs and add these "pseudo-GT" samples to the training dictionary.[16]

Parameter	**Recommended Setting for RTX 5070 Ti**	**Reasoning**
Input Resolution	112x112 or 160x160	Balances batch size with detail [28]
Max Frames	96 - 128	Critical for fast signing in DGS [15]
Batch Size	2 (Accumulate to 32)	Stable gradients for CTC/CE [29]
Learning Rate	$1 \times$ (Phase 1) $\rightarrow 5 \times$ (Phase 2)	Slow decay for Transformer stability [29]
VAC Weight	0.1 (Max)	Avoids CTC collapse [10]
Label Smoothing	0.1	Regularization for overconfident Attention [5]

Conclusions

The optimization of Continuous Sign Language Recognition on the PHOENIX-2014-T dataset is a delicate multi-objective problem. The gap between a 50% WER and the state-of-the-art is typically bridged by improving the model's awareness of motion and ensuring that the visual feature extractor receives granular, consistent supervision.[2] By moving away from resource-intensive 3D CNNs and instead utilizing Temporal Shift Modules, Adaptive Temporal Pooling, and Distillation-based CTC losses, it is possible to achieve competitive results on consumer-grade hardware.[6] The most significant improvements in 2024-2025 come from "sign-aware" pre-training and the use of learnable temporal filters, which allow the model to focus on the semantically rich portions of the video while ignoring the redundant background and transitional frames.[19] Implementing these techniques within a phased training regime—stabilizing alignment before introducing global context—provides the most robust pathway to high-accuracy recognition in the German Sign Language domain.[11]

Works cited

1. (PDF) A Comparative Study of Continuous Sign Language Recognition Techniques, accessed on March 19, 2026, https://www.researchgate.net/publication/381517846_A_Comparative_Study_of_Continuous_Sign_Language_Recognition_Techniques
2. Continuous Sign Language Recognition using CNN ... - Informatica, accessed on March 19, 2026, https://puffbird.ijs.si/index.php/informatica/article/view/8403/4474
3. Continual Learning in Multilingual Sign Language Translation - ACL Anthology, accessed on March 19, 2026, https://aclanthology.org/2025.naacl-long.546.pdf
4. RWTH-PHOENIX-Weather 2014 T: Parallel Corpus of Sign Language Video, Gloss and Translation, accessed on March 19, 2026, https://www-i6.informatik.rwth-aachen.de/~koller/RWTH-PHOENIX-2014-T/
5. Continuous Sign Language Recognition with Adapted Conformer via Unsupervised Pretraining - arXiv, accessed on March 19, 2026, https://arxiv.org/html/2405.12018v1
6. CorrNet+: Sign Language Recognition and Translation via Spatial-Temporal Correlation - arXiv, accessed on March 19, 2026, https://arxiv.org/html/2404.11111v1
7. Improving Continuous Sign Language Recognition with Cross-Lingual Signs - CVF Open Access, accessed on March 19, 2026, https://openaccess.thecvf.com/content/ICCV2023/papers/Wei_Improving_Continuous_Sign_Language_Recognition_with_Cross-Lingual_Signs_ICCV_2023_paper.pdf
8. [PDF] Continuous Sign Language Recognition Based on Cross ..., accessed on March 19, 2026, https://www.semanticscholar.org/paper/Continuous-sign-language-recognition-based-on-Zhu-Li/539c2d729819d561e0e9a157bb38c611ec6817fb
9. Li and Wang - 2020 - Reinterpreting CTC Training As Iterative Fitting - Scribd,

accessed on March 19, 2026, https://www.scribd.com/document/946953929/Li-and-Wang-2020-Reinterpreting-CTC-Training-as-Iterative-Fitting
10. Auxiliary Time-Alignment Losses in Neural Models - Emergent Mind, accessed on March 19, 2026, https://www.emergentmind.com/topics/auxiliary-time-alignment-losses
11. Towards Online Sign Language Recognition and Translation - arXiv.org, accessed on March 19, 2026, https://arxiv.org/html/2401.05336v1
12. Self-distillation Regularized Connectionist Temporal Classification Loss for Text Recognition: A Simple Yet Effective Approach - arXiv, accessed on March 19, 2026, https://arxiv.org/html/2308.08806v4
13. Self-Distillation Regularized Connectionist Temporal Classification Loss for Text Recognition: A Simple Yet Effective Approach - AAAI Publications, accessed on March 19, 2026, https://ojs.aaai.org/index.php/AAAI/article/view/28575/29118
14. 201224steven/Continuous-Sign-Language-Recognition-Based-on-Diffusion - GitHub, accessed on March 19, 2026, https://github.com/201224steven/Continuous-Sign-Language-Recognition-Based-on-Diffusion/
15. A Gloss-free Sign Language Production with Discrete Representation - arXiv, accessed on March 19, 2026, https://arxiv.org/html/2309.12179v2
16. Towards Online Continuous Sign Language Recognition and Translation - Scribd, accessed on March 19, 2026, https://www.scribd.com/document/932064187/Towards-Online-Continuous-Sign-Language-Recognition-and-Translation
17. Efficient Continuous Sign Language Recognition with Temporal Shift and Channel Attention, accessed on March 19, 2026, https://www.researchgate.net/publication/384765366_Efficient_Continuous_Sign_Language_Recognition_with_Temporal_Shift_and_Channel_Attention
18. Continuous Sign Language Recognition with Correlation Network - ResearchGate, accessed on March 19, 2026, https://www.researchgate.net/publication/373323848_Continuous_Sign_Language_Recognition_with_Correlation_Network
19. Dual-stage temporal perception network for continuous sign language recognition | Request PDF - ResearchGate, accessed on March 19, 2026, https://www.researchgate.net/publication/381281409_Dual-stage_temporal_perception_network_for_continuous_sign_language_recognition
20. Phoenix14t - CatalyzeX, accessed on March 19, 2026, https://www.catalyzex.com/s/Phoenix14t
21. SIGNKD: MULTI-MODAL HIERARCHICAL KNOWL- EDGE DISTILLATION FOR CONTINUOUS SIGN LAN- GUAGE RECOGNITION - OpenReview, accessed on March 19, 2026, https://openreview.net/pdf/7285b55cf96955ba29af6626d648a2141c44e118.pdf
22. CVPR Poster Continuous Sign Language Recognition With Correlation Network, accessed on March 19, 2026, https://cvpr.thecvf.com/virtual/2023/poster/23195
23. SignBERT+: Hand-model-aware Self-supervised Pre ... - arXiv.org, accessed on

March 19, 2026, https://arxiv.org/pdf/2305.04868
24. SignRep: Enhancing Self-Supervised Sign ... - CVF Open Access, accessed on March 19, 2026, https://openaccess.thecvf.com/content/ICCV2025/papers/Wong_SignRep_Enhancing_Self-Supervised_Sign_Representations_ICCV_2025_paper.pdf
25. SignBERT+: Hand-Model-Aware Self-Supervised Pre-Training for Sign Language Understanding | Request PDF - ResearchGate, accessed on March 19, 2026, https://www.researchgate.net/publication/370299522_SignBERT_Hand-model-aware_Self-supervised_Pre-training_for_Sign_Language_Understanding
26. SignBERT+: Hand-Model-Aware Self-Supervised Pre-Training for Sign Language Understanding - PubMed, accessed on March 19, 2026, https://pubmed.ncbi.nlm.nih.gov/37099464/
27. Continuous Sign Language Recognition Based on Cross-Resolution Knowledge Distillation, accessed on March 19, 2026, https://www.springerprofessional.de/en/continuous-sign-language-recognition-based-on-cross-resolution-k/27469478
28. [2303.06820] Continuous sign language recognition based on cross-resolution knowledge distillation - arXiv.org, accessed on March 19, 2026, https://arxiv.org/abs/2303.06820
29. Track: Poster Session 6 - NeurIPS, accessed on March 19, 2026, https://neurips.cc/virtual/2022/session/64307